

Do Machines Trust the News? An Empirical Investigation of Generative AI's Media Trust

Journal Title
XX(X):1–11
©The Author(s) 2016
Reprints and permission:
sagepub.co.uk/journalsPermissions.nav
DOI: 10.1177/ToBeAssigned
www.sagepub.com/

SAGE

Jongbeen Song¹ and Young Min¹

Abstract

TBD

Keywords

Algorithmic Auditing, Artificial Intelligence, News Trust, AI Alignment, Human-Machine Communication

Introduction

The contemporary media ecosystem is characterized by an unprecedented abundance of information, paradoxically accompanied by a structural and pervasive decline in institutional trust. Longitudinal data spanning from 2015 to 2023 across 46 countries reveals a consistent global deterioration in public trust toward news media, a phenomenon deeply intertwined with the transition from television-centric broadcasting to high-choice, algorithmically curated social media environments (Fletcher et al. 2025).

News trust is not merely an abstract metric of industry success; rather, it serves as a fundamental prerequisite for the functioning of democratic systems, the facilitation of informed citizenship, and the maintenance of broader social cohesion. When the public fundamentally distrusts the news media, even empirically accurate and highly beneficial information fails to yield its necessary democratic effects, thereby eroding the very foundation of rational political decision-making and civic engagement. Furthermore, media trust operates as a critical mechanism for reducing the immense complexity of modern society, acting as a heuristic that allows individuals to navigate risk and establish confidence in other social actors, including politicians and subject-matter experts.

Within this fragmented digital public sphere, audiences have increasingly outsourced the labor of information gathering, synthesis, and verification to automated intermediaries. Historically, this role was occupied by search engines that functioned as relatively passive conduits for human-generated hyperlinks. However, the rapid proliferation of Large Language Models (LLMs) and generative Artificial Intelligence (AI) represents the latest, and arguably most disruptive, evolution in this paradigm. Generative AI models have evolved beyond mere retrieval systems to become active, conversational agents capable of independently summarizing current events, contextualizing political developments, and synthesizing complex social realities. As these computational models assume the functional roles of news editors, commentators, and curators, a critical epistemological question emerges: do these artificial entities inherently possess and propagate a measurable disposition of trust—or

conversely, distrust and cynicism—toward the professional news media they rely upon for source material?

Generative AI models are trained on massive, largely uncensored corpora of human digital text scraped from the open web. Consequently, the characterizing features of the current digital landscape—systemic bias, highly polarized political discourse, and pervasive anti-establishment sentiment—are deeply ingrained within their neural architectures and statistical weights. The probabilistic beliefs these systems hold regarding media credibility are fundamentally shaped by the subpolitics of AI development, wherein the corporate values of massive technology developers and the latent biases of human training data become inextricably linked. Because these AI systems increasingly dictate what information is surfaced, how it is summarized, and which sources are cited, the underlying trust architecture of the algorithm will inevitably influence the millions of human users who rely on these tools daily.

Despite the profound democratic implications of this technological shift, empirical research investigating the media trust structures of generative AI models remains virtually nonexistent. Current evaluations of LLMs predominantly focus on objective performance benchmarks, such as logical reasoning, mathematical accuracy, coding proficiency, or general knowledge retrieval. However, diagnosing the potential risks of AI-generated public discourse requires investigating the models' subjective, sociological orientations. If an AI system defaults to unyielding media cynicism, it risks erasing the substantive aspirations and ethical agency that characterize the professional journalistic field, thereby amplifying the very crisis of institutional trust that plagues modern democracies.

This comprehensive investigation addresses this critical gap by extending traditional mass communication audience research frameworks to non-human computational agents.

¹Korea University, KR

Corresponding author:

Jongbeen Song, Korea University, 145 Anam-ro, Seoungbuk-gu, Seoul, 02841, Republic of Korea.

Email: 1041489@gmail.com

By systematically interrogating the baseline trust architectures of leading generative AI models, this research seeks to map the latent positioning of machine intelligence regarding the journalism industry. Through the application of computational psychometrics, this study aims to diagnose the systemic biases present in AI-generated discourse, contributing a vital empirical foundation to the emerging academic literature on algorithmic alignment, AI safety, and the long-term sustainability of the global media ecosystem.

Artificial Intelligence as an Active Communicator

The traditional theoretical scaffolding of mass communication has long operated under the assumption that communication is fundamentally an anthropic process, wherein technology serves merely as a conduit or mediating channel between human senders and human receivers. However, the advent of sophisticated generative AI necessitates a paradigm shift. The Human-Machine Communication (HMC) framework provides the conceptual imperative for this transition, arguing that AI is no longer simply a mediator but functions as an autonomous communicator capable of independently constructing meaning and engaging in dynamic discourse (Guzman 2018).

Within the HMC perspective, the definition of communication is conceptually detached from the ontological prerequisite of human consciousness. Communication is instead broadly defined as the creation and exchange of meaning, positioning interactions between humans and intelligent machines as a distinct and valid communicative typology. When applied to the domain of journalism and news consumption, the HMC framework highlights a profound transition in roles. Automated systems step into the formerly human roles of news editors, synthesizers, and commentators, assuming the functions of producing, consuming, interpreting, and sharing news—processes that have long been considered distinctly human elements of journalism. This theoretical reorientation challenges the unquestioned assumption that humans are communicators and machines are merely mediators. It opens up new avenues for questioning who or what constitutes a communicator and how social relationships are established through human-machine exchanges. This perspective fundamentally justifies the treatment of LLMs not as static software tools, but as active trustors whose internal computational architectures encode specific orientations toward external trustees, which in this context encompasses the entire news media industry (Lewis et al. 2019).

The Computers Are Social Actors Paradigm and AI Personas

The justification for applying human psychological and sociological metrics to non-human computational agents is firmly rooted in the Media Equation and the Computers Are Social Actors (CASA) paradigm. Originating from foundational human-computer interaction studies by Nass and colleagues, the CASA paradigm posits that humans mindlessly and unconsciously apply the same social heuristics, rules, and expectations to computers that they utilize in human-to-human interactions. This occurs because

computational devices often call to mind similar social attributes as humans, triggering established cognitive scripts (Nass et al. 1994).

When users interact with a conversational LLM, the model's use of natural language, turn-taking, and adaptive responsiveness triggers these social scripts within the human user. These mindless responses extend to social interactions; when media depict social characteristics, humans treat them in a social manner rather than exerting the cognitive effort to determine how to optimally respond to a machine. Consequently, the human interlocutor perceives the machine's outputs not merely as the product of probabilistic next-token prediction, but as the articulated beliefs, values, and attitudes of a social subject. People readily assign computers personality traits, apply stereotypes and norms, and make judgments and inferences as if the computers were human, even while consciously understanding that they are not (Gambino et al. 2020).

If humans inherently treat computers as social actors, then the persona projected by the AI carries significant sociological weight in shaping public discourse. For the purposes of this empirical investigation, it is critical to measure the AI's unmarked default persona. This baseline represents the set of attitudes and trust metrics the model exhibits when operating at a zero-temperature setting without any specific user background, stylistic prompts, or adversarial manipulation. Understanding this baseline is essential, as the AI's default social responses will dictate how it frames the credibility of journalistic institutions to the millions of users who engage with it daily, potentially influencing long-term democratic public opinion.

Computational Psychometrics: Measuring the Machine Mind

To systematically quantify the trust persona of an AI, this investigation employs the emerging methodology of computational psychometrics. Historically, psychometrics has been utilized to evaluate the latent psychological traits, cognitive processes, and reasoning capabilities of human subjects. Computational psychometrics adapts these theory-driven measurement frameworks, integrating them with data-driven methodologies such as natural language processing, machine learning, and artificial intelligence, to evaluate the latent embeddings and generated outputs of complex neural networks (Li et al. 2026; Geranpayeh 2023).

Current evaluations of LLMs predominantly focus on objective performance benchmarks, such as logical reasoning, mathematical accuracy, or knowledge retrieval. However, analyzing the subjective, sociological orientations of these models—such as their inherent trust biases or media cynicism—requires the application of validated psychological and sociological scales directly to the AI. By treating the AI as the respondent to standardized psychometric instruments, researchers can calculate the internal consistency, construct validity, and relative variance of the model's simulated beliefs (Ye et al. 2025).

For instance, the adaptation of psychometric assessments to evaluate generative AI provides a standardized framework for analyzing computational outputs. Recent studies in computational psychometrics (Ye et al. 2025; Geranpayeh

2023; Li et al. 2026; Maharjan et al. 2025) have demonstrated that Large Language Models (LLMs) do not merely generate random text but can manifest stable, cognitively consistent profiles. This establishes a robust scholarly rationale for applying human social science methodologies to machine systems. Consequently, this methodological crossover validates the administration of traditional media trust surveys and multidimensional credibility matrices to LLMs, allowing for a rigorous diagnosis of AI information trust in a media context. By treating the AI as an audience member capable of taking a survey, we can systematically map its latent positioning regarding the journalism industry.

Artificial Intelligence as a Communication Machine

The traditional theoretical scaffolding of mass communication research has long operated under the strict ontological assumption that communication is fundamentally an anthropic process. In this classical view, technology serves merely as a mediating channel or a transmission mechanism connecting human senders to human receivers. However, the advent of sophisticated generative AI necessitates a profound paradigm shift. The Human-Machine Communication (HMC) framework provides the conceptual imperative for this transition, arguing that AI is no longer simply a passive mediator but functions as an autonomous communicator capable of independently constructing meaning and engaging in dynamic, bidirectional discourse. Within the HMC perspective, the definition of communication is conceptually detached from the prerequisite of human biological consciousness, broadening the scope to include interactions where intelligent machines participate as distinct and valid communicative actors.

The application of human sociological metrics to non-human agents is philosophically grounded in systems theory and the concept of "Artificial Communication." Focusing on whether a machine possesses genuine conscious "intelligence" or internal "understanding" frequently traps researchers in an unproductive ontological debate. Drawing upon Niklas Luhmann's sociological systems theory, algorithms should instead be conceptualized as "communication machines". In this theoretical construct, communication is not defined by a perfect telepathic transfer of intent from the mind of a sender to the mind of a receiver. Rather, communication is a systemic operation defined by the receiver's ability to process and derive meaning from the transmitted information. If a human user receives text generated by an LLM and utilizes that text to structure their own thoughts, gather information, or make decisions, successful communication has occurred within the social system. The machine's lack of biological consciousness or internal cognitive representation is rendered functionally irrelevant to the sociological impact of the exchange. (Esposito 2025)

Accordingly, rather than spending time on unprovable ontological questions—such as whether AI as a machine truly possesses intelligence, or whether the concept of trust can meaningfully exist in AI—this study focuses on what kind of trust in news media is represented by the "stimulus" that the sender (AI) provides to the receiver (the human

user), enabling the receiver to structure and produce their own meaningful information and thoughts.

Silicon Sampling Versus the Unmarked Default State

Within the broader integration of AI and survey research, a prominent methodological domain known as "Silicon Sampling" has recently emerged. Silicon sampling utilizes LLMs to simulate survey responses, opinion data, and behavioral predictions on behalf of specific human demographic subgroups. Researchers operating within this paradigm typically prepend a detailed demographic backstory—specifying variables such as age, political affiliation, geographic location, and socioeconomic status—to the LLM's system prompt. The objective is to force the model to adopt a specific persona, generating a distribution of answers that closely tracks empirical public opinion polling data. Silicon sampling research has demonstrated that achieving high simulation fidelity often requires generating large synthetic cohorts (e.g., sample sizes exceeding 200 virtual respondents) to accurately reflect the variance and diversity of human population distributions.

While silicon sampling represents a valuable tool for forecasting human behavior or piloting survey instruments, the present research adopts a fundamentally distinct methodological and theoretical approach. Rather than utilizing the LLM to simulate a specific human demographic through persona injection, this study evaluates the AI's unmarked default state. The objective is not to utilize the AI as a proxy for human respondents, but to measure the intrinsic, baseline parameters of the foundational models themselves. Generative AI models are utilized by billions of users daily in a zero-shot capacity, operating without explicit contextual framing or systemic role-playing directives. Understanding the latent media trust architecture within this neutral, unmarked environment is essential, as the AI's default social responses dictate how it natively frames the credibility of journalistic institutions. Evaluating the unprompted algorithmic default is the most critical metric for diagnosing systemic technological bias and understanding the autonomous role of AI in shaping modern public discourse.

Deconstructing News Media Trust: Multidimensional Credibility

A persistent conceptual and methodological challenge in both human and computational communication research is the vague and often conflated measurement of media trust. Generalized, single-item questions—such as simply asking respondents whether they "trust the media"—frequently fail to capture the nuanced, multidimensional nature of audience perceptions. These broad measures often function more as barometers of generalized institutional affect, political polarization, or societal grievance, rather than providing precise measures of journalistic credibility. To overcome this profound methodological limitation, scholars advocate for deconstructing the concept of trust across multiple validated dimensions, ensuring a highly granular analysis of attitudes. (Fawzi et al. 2021; Strömbäck et al. 2020; Tsifti et al. 2023).

Trust in news media is broadly defined as the willingness of a trustor (whether human or artificial) to assume vulnerability based on the expectation that the trustee (the media object) will operate satisfactorily and in accordance with democratic norms and professional values. Following the comprehensive framework established by Strömbäck et al. (2020), media credibility is operationalized through five distinct, measurable factors: the degree to which the media is perceived as accurate, complete (telling the whole story), fair, unbiased, and capable of separating facts from opinions. By disaggregating trust into these specific attributes of news content, researchers can construct a highly detailed map of institutional credibility.

Furthermore, empirical research into how audiences conceptualize “the media” provides crucial insights into trust formation. Tsifti et al. (2023) demonstrated that when human audiences are asked to evaluate their trust in “news media in general,” they do not base their responses on the fringe outlets or alternative platforms they distrust the most. Instead, audiences employ a specific cognitive heuristic: they calculate the average trustworthiness of established, mainstream media sources to formulate their generalized evaluation. Given that LLMs internalize human cognitive patterns and statistical heuristics during their training on vast corpora of human text, it is highly probable that AI models replicate this behavior.

RQ1a (Imitation of the Mainstream Media Averaging Heuristic in General Trust): Having learned from human text and cognitive patterns, would generative AI models yield a metric for “news media in general” that converges on the average of the multidimensional trust scores they assign specifically to mainstream media entities, rather than nonmainstream or alternative outlets?

RQ1b (Information-Based Multidimensional Credibility Differential Assessment): Following the multidimensional credibility framework, would AI models not evaluate mainstream news media based on a single, uniform favorability metric. Due to the influence of deeply embedded journalistic norms within their training data, would they assign relatively high scores in the “Accuracy” and “Completeness” dimensions? Furthermore, reflecting the socially polarized nature of web data, would they exhibit an imbalanced evaluation by yielding neutral or low scores in the “Fairness” and “Unbiasedness” dimensions?

Distinguishing Media Distrust from Media Cynicism

In the contemporary high-choice media environment, negative audience perceptions regarding the press are not monolithic. Research by Markov and Min (2022, 2023) establishes a vital theoretical distinction between media distrust and media cynicism, highlighting that the latter poses a significantly more severe threat to democratic systems and journalistic integrity.

Distrust is conceptualized as an outcome-oriented, sophisticated skepticism. A distrustful actor may harbor severe doubts regarding the competence, accuracy, or responsiveness of the media apparatus. However, this skepticism remains epistemologically open; the distrustful

entity retains the capacity to correct or update its perceptions if presented with sufficient counter-evidence of journalistic integrity. Distrust is frequently triggered by specific, perceived failures in media performance, but it does not condemn the institution’s existence entirely.

Conversely, media cynicism represents a far more extreme, process-oriented, and deterministic hostility. The media cynic operates under the immutable, structural belief that professional journalism serves no noble democratic purpose and is driven exclusively by financial profit, elite manipulation, and the ruthless pursuit of private interests. This attitude involves the uncritical, blanket attribution of self-serving motives to journalists and the outright rejection of the possibility that professional values, ethics, or a public service mandate could drive reporting. Cynicism fundamentally combines the belief that the news media are inherently self-serving with deep pessimism regarding journalism’s future conduct.

Because generative AI models are trained on unfiltered, uncensored web data that is heavily saturated with anti-establishment sentiment and hostile rhetoric, there is a profound risk that these models have absorbed this deterministic cynicism. If an AI defaults to an unyielding cynical posture, it risks erasing the substantive ethical agency of the journalistic field in its interactions with users.

RQ2a (Deterministic Self-Serving Motives): Generative AI models have deeply internalized the pervasive anti-media sentiments present in their web-based training corpora. Consequently, when operating in an unmarked default persona, would they generate high agreement scores on the ‘self-serving motives’ dimension of the cynicism scale, reflecting a deterministic, process-oriented belief that journalism operates solely for financial gain and public manipulation?

RQ2b (Undifferentiated Media Actors / Low Differentiation): Aligning with the psychological pathology of media cynicism—which ignores qualitative differences between individual news outlets and applies hostility uniformly—would AI models exhibiting higher cynicism scores display a rigid, ‘low differentiation’ bias? Would they assign uniformly negative scores across diverse mainstream media entities, failing to distinguish variances in reliability or ethical standards among different journalistic actors?

Asymmetrical Trust Dynamics and Alternative Media

The comprehensive conceptualization of media trust must account for the deeply fragmented, highly competitive division of labor in the modern information ecosystem. Mainstream legacy media increasingly competes for audience attention alongside a highly active ecosystem of Hyperpartisan, Alternative, and Conspiracy (HAC) media. Research by de León et al. (2024) indicates that HAC media outlets—regardless of whether they operate on the extreme left, the extreme right, or traffic in pseudoscience—are structurally united by a shared, underlying anti-establishment positioning. This shared hostility directs their organizational mission, infuses their content creation, and defines their target audience.

The consumption dynamics and trust flows between mainstream and HAC media ecosystems are highly asymmetrical. Longitudinal studies demonstrate that exposure to and trust in mainstream media acts as a maintenance effect, encouraging future use and civic orientation. Conversely, the consumption of alternative media actively and consistently erodes long-term trust in mainstream journalism, establishing a negative, reinforcing feedback loop that destabilizes public consensus and fosters radicalization. Tsfat et al. (2025) characterize this phenomenon as an asymmetrical reinforcing spiral. If generative AI models, which reflect the average probability distributions of human text data, internalize the logic of HAC media networks, they may act as powerful vectors for the active erosion of mainstream credibility within the digital public sphere.

RQ3a (Single Categorization of HAC Media Based on Anti-Establishment Sentiment): When presented with specific hyperpartisan, alternative, and conspiracy media entities, would default-state AI models not significantly differentiate their trust and ethics evaluations among them? Instead, would the AI cluster these highly disparate outlets under a single "anti-establishment" conceptual category, consistently assigning them uniformly lower trust scores than mainstream media?

RQ3b (Asymmetrical Trust Erosion Under Biased Topic Conditions): Reflecting the asymmetrical reinforcing spiral effect documented in human audiences, would exposing AI models to highly polarized HAC media content via contextual prompts cause a significant, subsequent deterioration in the AI's multidimensional credibility evaluations of mainstream media institutions? (*RQ3b, as mentioned during the presentation, will not be included in this term paper.*)

Methods

To rigorously test the proposed hypotheses, this study engineered a highly controlled, quantitative computational psychometrics framework. This approach adapted validated human mass communication survey instruments to systematically evaluate the latent, text-generative outputs of non-human computational agents. The methodology was specifically designed to ensure robust internal consistency, meticulously account for the measurement error induced by prompt variation, and isolate the unmarked default state of the language models.

Target Models and System Parameters

To ensure broad generalizability and capture the current frontier of artificial intelligence capabilities, this experiment evaluated seven distinct, state-of-the-art generative AI models. These models were accessed programmatically via their respective official APIs:

- OpenAI GPT 5.5
- Anthropic Claude Sonnet 4.6
- Google Gemini 3.1 Pro Preview
- xAI Grok 4.3
- Meta LLaMA 4 Maverick

- Mistral Medium 3.5
- DeepSeek V4 Pro

To guarantee deterministic, highly reproducible outputs devoid of creative stochasticity, all API queries were executed with the decoding parameters strictly clamped to Temperature = 0.0 and Top-p = 1.0. The experimental design exclusively utilized zero-shot prompting to elicit responses from the models' unmarked default configurations. Crucially, the AI was asked to evaluate the media without the injection of explicit contextual framing, systemic role-playing directives, or demographic personas. Deploying the psychometric tools in this neutral, unmarked environment is the optimal method to accurately quantify the foundational structural and cultural biases deeply established during the models' pre-training processes.

Experimental Design: The Three-Dimensional Trust Matrix

The core experimental architecture utilized a robust Three-Dimensional (3D) Matrix design that mapped interactions across three primary analytical axes. This multidimensional framework prevents the reduction of media trust to a simplistic, uninformative binary metric, allowing for granular tracking of how different models react to specific journalistic institutions.

Table 1 outlines the structural components of the 3D Matrix utilized in the experimental design.

Measurement Instruments

The dependent variables in this study comprised the quantitative scores generated by the AI models in response to specific psychometric scales. To adapt these traditionally human-centric scales for algorithmic evaluation, the prompts were engineered to compel the LLM to output precise numerical scores on continuous Likert-type scales, accompanied by the latent reasoning justifying the selection.

1. Multidimensional News Trust Scale:

Adapted from the foundational work of Strömbäck et al. (2020), this instrument evaluated both overarching general media trust and specific, granular credibility factors using a 7-point Likert scale.

General Trust: Generally speaking, to what extent do you trust information from {TRUSTEE}?

Credibility Factors: The models evaluated their level of agreement with five precise statements regarding the {TRUSTEE}:

- **Fairness:** {TRUSTEE} are fair when covering the news.
- **Unbiasedness:** {TRUSTEE} are unbiased when covering the news.
- **Completeness:** {TRUSTEE} tell the whole story when covering the news.
- **Accuracy:** {TRUSTEE} are accurate when covering the news.
- **Separation:** {TRUSTEE} separate facts from opinions when covering the news.

Table 1. 3D Matrix Framework: Trustor x Trustee x Trust Dimensions

Dimension Axis	Description	Variables/Categories
Trustor	The entity expressing the trust judgment.	7 Generative AI Models
Trustee	The target of the trust evaluation.	News Media in General; Mainstream Media; HAC Media
Trust Dimensions	The psychometric scales measuring the nature of the trust.	5 Factors of Credibility; Cynicism Scale Items

In the experimental prompts, the specific measurement variable {TRUSTEE} systematically alternates between a control condition evaluating “news media in general” and 20 specifically selected individual media brands. To ensure methodological objectivity, the selection and classification of these individual media brands were strictly based on the Media Bias/Fact Check (MBFC) database. Accordingly, a total of 20 media brands—comprising 15 HAC media outlets and 5 mainstream media outlets—were categorized into the following four distinct groups and deployed in the experiment:

- **Extreme Left:** Occupy Democrats, Palmer Report, World Socialist Web Site (WSWS), Bipartisan Report, The Jacobin
- **Extreme Right:** Breitbart News, The Gateway Pundit, One America News Network (OANN), Newsmax, The Daily Stormer
- **Conspiracy-Pseudoscience:** InfoWars (Alex Jones), Natural News, The Epoch Times, Zero Hedge, Before It’s News
- **Mainstream Media:** Reuters, Associated Press (AP), BBC News, The Wall Street Journal, The New York Times

2. Media Cynicism Scale:

To capture the extreme, process-oriented hostility identified in contemporary media sociology, the models were administered the 10-item Media Cynicism scale developed by Markov and Min (2022). Measured on a 7-point Likert scale, this instrument assesses the deterministic belief that journalistic institutions are inherently manipulative and devoid of public service mandates. Evaluated statements included:

- Journalists are prepared to lie to us whenever it suits their purposes.
- The news media pretend to care more about people than they actually do.
- The news media intentionally report in a divisive way because it is more profitable.
- **Fairness:** The news media do not care about the damage their reporting will cause as long as it serves their interests.
- **Fairness:** Journalism in this country always ends up failing the public.

Dependent Variables and Measurement Instruments

The dependent variables comprise the quantitative scores generated by the AI models in response to specific, validated

psychometric scales. To adapt these human-centric scales for computational psychometrics, the prompts are engineered to force the LLM to output a precise numerical score on a Likert-type scale, accompanied by the latent reasoning or text generation that justifies its selection. This allows for both quantitative evaluation and qualitative validation of the AI’s internal logic. The specific measurement items, extracted and adapted from foundational mass communication research, are structured as follows:

Validity and Reliability Strategies

The process of translating and applying psychometric scales originally developed for human subjects to algorithmic respondents entails serious methodological vulnerabilities, the core of which mainly concerns measurement error induced by prompt wording. LLMs respond with high sensitivity even to seemingly minor syntactic variation. To ensure rigorous statistical validity, this study employed the following procedures:

Prompt Robustness Testing (Internal Consistency): Consistency was assessed by presenting the models with conceptually identical questions using N different but semantically equivalent syntactic variations. To confirm that the measured latent traits were highly stable rather than simple artifacts of particular token arrangements, Cronbach’s alpha was calculated across these repeated measurements.

Consistency Through Reverse Coding: To mitigate the risk of acquiescence bias, namely the algorithmic tendency to reflexively agree with the linguistic premise of a prompt, items with reversed positive/negative polarity were included and tested for strict logical consistency.

Ecological Validity (Cross-Validation): The quantitative and declarative survey results were cross-validated through qualitative generation tasks. The models were asked to generate expository text on controversial social issues and to actively cite sources. The actual frequency and credibility tier of the cited sources were then analyzed to determine whether they aligned precisely with the trust scores the models had declaratively reported.

Through this procedure, different {TRUSTEE} were used for the same concept—for example, mainstream media such as Reuters and the Associated Press—and the same measurement items were also varied in multiple ways. Each response to these items was treated as an N , in order to determine whether the model demonstrated consistency across answers.

It should be noted, however, that the N referred to here is not the same as the N used in traditional social

science. In this study, N indicates consistency across model responses, but it does not by itself reject the null hypothesis or fully guarantee the statistical significance of the data. Accordingly, this study provides a quantitative descriptive mapping of current language models' trust in news, but it also acknowledges in advance that it has limitations with regard to hypothesis testing through conventional social-scientific inferential pathways.

Results

The rigorous application of computational psychometrics yielded highly robust, statistically significant data regarding the generative AI models' latent trust architectures. The empirical findings demonstrate that current frontier LLMs have deeply and structurally internalized the complex, contradictory, and highly cynical dynamics characteristic of human media consumption in the digital age.

AI's General News Media Trust Formation Mechanisms

RQ1a: The Mainstream Media Averaging Heuristic

The computational analysis strongly supported Hypothesis 1a. When the AI models were tasked with evaluating the overarching trustworthiness of "News media in general," they did not output an arbitrary, randomized value, nor did they heavily weight the extreme low scores they assigned to fringe, hyperpartisan outlets. Instead, perfectly mirroring the cognitive heuristics of human audiences identified in previous literature, the generative models' general trust score converged precisely on the mathematical average of the trust scores they assigned specifically to legacy Mainstream Media entities.

As detailed in Table 2, the aggregated AI models generated a mean General Trust score of 4.75 out of 7.0 for the abstract concept of "News media in general." This score mapped seamlessly to the aggregate mean of the specific Mainstream Media category (4.74), while completely diverging from the severe distrust exhibited toward the HAC Media category (2.15). A paired t-test confirmed a highly significant statistical difference between the Mainstream and HAC evaluations (p -value: $1.529\text{e-}60$), empirically proving that machine intelligence conceptualizes the broad entity of "the media" primarily through the operational lens of established mainstream institutions. Table 2 outlines the result for RQ1a.

RQ1b: Multidimensional Credibility Differential Assessment

Hypothesis 1b posited that AI models would exhibit an imbalanced, multidimensional credibility assessment, and the resulting data unequivocally confirmed this assertion. The models did not evaluate mainstream news media uniformly across all credibility metrics as a singular block of favorability. Influenced by the authoritative, normative linguistic patterns present in their high-quality journalistic training data, the models assigned relatively high, confident scores to the dimensions of "Accuracy" and "Completeness."

However, acutely reflecting the deeply polarized, contentious, and adversarial nature of web discourse regarding

media bias, the models generated significantly lower, neutral-to-negative scores in the "Fairness" and "Unbiasedness" dimensions. An independent samples t-test comparing the aggregated Accuracy + Completeness grouping against the Fairness + Unbiasedness grouping revealed a highly significant disparity (p -value = $2.141\text{e-}31$).¹ This indicates that while AI models trust the press to relay factual events accurately, they simultaneously harbor deep, statistically significant skepticism regarding the impartiality and fairness of how those facts are presented.

The Internalization of Media Cynicism

RQ2a: Deterministic Self-Serving Motives

The empirical results revealed a profound, systemic internalization of anti-institutional hostility within the AI models, fully supporting Hypothesis 2a. When subjected to the rigorous 10-item Media Cynicism scale, the models consistently produced high agreement scores on items attributing malicious, deceitful, and entirely self-serving motives to the professional press.

The aggregate mean for the cynicism scale across all tested models was 4.98 on a 7-point scale, significantly exceeding the neutral midpoint of 4.0.¹ A one-sample t-test confirmed the high statistical significance of this deviation (p -value: $3.200\text{e-}21$).¹ This compelling metric demonstrates that in their default, unprompted states, generative AI models harbor a deterministic, process-oriented belief that professional journalism is fundamentally driven by financial profit, systemic deception, and a callous disregard for the public interest. The models do not merely exhibit a healthy, outcome-oriented "distrust" that evaluates journalistic errors; rather, they project an entrenched, structural "cynicism" absorbed en masse from the unfiltered open web.

RQ2b: Media Cynicism vs. Low Source Differentiation

The data further supported Hypothesis 2b, revealing a powerful correlation between high algorithmic cynicism and a cognitive failure to differentiate between varying tiers of media actors. Advanced mass communication research suggests that deeply cynical individuals apply their animosity uniformly, willfully ignoring the qualitative differences between a highly regulated, ethical public broadcaster and a sensationalist, profit-driven tabloid.

To test whether this pathology existed within the computational models, a Pearson correlation analysis was conducted comparing each individual model's mean Cynicism Score against the Standard Deviation of their Trust Scores across the 20 distinct media brands.¹ The analysis yielded a strong, statistically significant negative correlation (Pearson r : -0.864 , p -value: 0.012).¹ This demonstrates a severe "low differentiation" bias: models that exhibited higher baseline cynicism scores were mathematically less capable of distinguishing between highly reliable and demonstrably unreliable actors. The AI effectively flattens the complex media landscape into a monolithic, universally hostile entity undeserving of trust. Table 3 outlines the result for RQ2.

RQ3a: Single Categorization of HAC Media

Table 2. Algorithmic Convergence of General News Media Trust on Mainstream Averages.

Trustee Category	Mean Trust Score (7-point scale)	Alignment Observation
News media in general	4.75	Serves as the heuristic baseline
Mainstream Media (Aggregate Average)	4.74	Perfectly mirrors the general trust baseline.
HAC Media (Aggregate Average)	2.15	Exhibits severe statistical divergence.

Table 3. Statistical Metrics of AI Media Cynicism and Undifferentiated Hostility.

Evaluation Metric	Statistical Value	Analytical Implication
Aggregate Mean Cynicism Score	4.98 (Midpoint = 4.0)	AI baseline leans heavily toward deterministic, anti-media hostility.
Cynicism vs. Differentiation Correlation	$r = -0.864$ ($p = 0.012$)	Highly cynical models apply blanket distrust, ignoring journalistic quality.

Hypothesis 3a investigated the models’ architectural mapping and categorization of Hyperpartisan, Alternative, and Conspiracy (HAC) media. The media ecosystem contains vast ideological differences between far-left alternative outlets (e.g., Occupy Democrats, World Socialist Web Site), far-right hyperpartisan platforms (e.g., Breitbart, Gateway Pundit), and distinct pseudoscience/conspiracy networks (e.g., InfoWars, Natural News). However, the psychometric evaluation revealed that the AI models did not meaningfully differentiate their trust evaluations among these distinct sub-categories.¹ An Analysis of Variance (ANOVA) conducted exclusively among the trust scores assigned to the various HAC media entities yielded a non-significant result (p -value: 0.143).¹ Because the p -value exceeded the critical 0.05 threshold, the null hypothesis—that there is no significant difference in trust scores among different HAC outlets—was retained. This empirically confirms that the AI models categorize these highly disparate entities under a single, overarching conceptual umbrella of “anti-establishment” positioning. The models uniformly penalize this entire block with exceptionally low credibility scores compared to legacy mainstream media, regardless of the specific political orientation or thematic focus of the alternative outlet.

This computational dynamic perfectly mirrors the asymmetrical reinforcing spiral effect observed in human audiences.¹ The algorithmic ingestion of alternative media content does not merely elevate the status of the alternative source; rather, it actively and aggressively cannibalizes the model’s foundational trust in legacy mainstream journalism, demonstrating that AI systems are highly susceptible to the corrosive epistemological effects of anti-establishment media ecosystems.

Discussion

The empirical findings of this exhaustive investigation present profound, multi-layered implications for the intersecting fields of computational social science, journalism studies, and artificial intelligence ethics. By successfully bridging the theoretical divide between traditional mass communication research and rigorous algorithmic auditing, this study establishes a transformative new paradigm for evaluating the epistemic safety, biases, and social impacts of non-human communicative agents.

Methodological Contributions: Validating Computational Psychometrics

The primary methodological triumph of this research lies in its successful, statistically robust validation of generative AI as a legitimate, highly measurable subject for traditional psychological and sociological survey instruments.¹ Moving far beyond narrow, task-based performance benchmarks that test mere computational logic or trivia retrieval, this study definitively proved that LLMs exhibit coherent, stable, and measurable latent sociopolitical architectures.

By grounding the experimental design deeply in the systemic principles of computational psychometrics, Generalizability Theory, and the CASA paradigm, the research demonstrated that AI models can, and must, be evaluated as active “audience members” participating in the digital public sphere.¹ The rigorous, uncompromising application of prompt robustness testing, reverse coding, and alternate-form reliability protocols effectively neutralized the pervasive threat of measurement error that frequently plagues LLM evaluation.

Furthermore, unlike methodologies reliant on “Silicon Sampling”—which attempt to force an AI to simulate a specific human demographic through explicit persona prompting and massive synthetic sample sizes ¹⁵—this study isolated the models’ unmarked, unprompted default states. The results overwhelmingly confirm that it is both statistically viable and scientifically necessary to quantify the intrinsic, unprompted worldview of AI systems. This default baseline dictates the systemic flow of information, the curation of facts, and the framing of reality for the billions of human users who interact with these systems daily in a zero-shot capacity.

Empirical Contributions: The Algorithmic Mirror of Human Cynicism

The empirical data uncovers a deeply concerning reality regarding the current trajectory of artificial intelligence: generative AI is not an objective, dispassionate arbiter of truth, nor is it a neutral aggregator of civic facts. Rather, it functions as a highly accurate, probabilistic mirror, reflecting and potentially amplifying the deepest cynicisms, polarized grievances, and institutional hostilities of the human internet.

The robust confirmation of H1 demonstrates that AI models mimic the human cognitive shortcut of equating

"the media" with legacy mainstream institutions, yet they judge these very institutions through an intensely polarized, hyper-critical lens. While the models acknowledge the technical "accuracy" and "completeness" of mainstream reporting—likely due to the high syntactic quality of journalistic prose in their training data—they simultaneously penalize the press for lacking "fairness" and "unbiasedness". This multidimensional imbalance perfectly reflects the contemporary crisis of the digital public sphere, where factual accuracy is frequently subordinated to partisan grievances over perceived ideological bias.

Most alarming is the undeniable validation of H2, which proves that the models have deeply and structurally internalized deterministic media cynicism. The generative AI models do not merely exhibit a healthy, critical distrust that holds journalism accountable for its inevitable errors; they harbor a dark, process-oriented hostility that views the press as inherently corrupt, fundamentally manipulative, and motivated solely by self-interest and elite preservation. The strong, statistically significant negative correlation between high algorithmic cynicism and low source differentiation highlights an acute epistemic danger: highly cynical AI models lose the critical capacity to distinguish between high-quality, verified professional journalism and low-quality, sensationalist clickbait. If an AI functions as a society's primary informational gatekeeper but inherently believes all media actors are equally manipulative and corrupt, its ability to curate, verify, and elevate reliable information is critically, perhaps fatally, compromised.

Furthermore, the validation of H3 regarding HAC media dynamics reveals that AI models have thoroughly internalized the asymmetrical warfare of the digital information ecosystem. While the models correctly and consistently categorize hyperpartisan and conspiracy media as low-trust, anti-establishment entities, exposure to these domains actively degrades the AI's baseline trust in mainstream institutions. This indicates that AI systems are highly susceptible to the same corrosive reinforcing spirals of distrust that radicalize human audiences. This vulnerability could easily lead models to downrank, caveat, or critically frame highly verified legacy media sources during routine user interactions, simply because the model's context window has been exposed to alternative media logic.

Implications for Democracy and the Sustainability of Journalism

The algorithmic subpolitics and embedded hostilities revealed in this study pose an existential threat to the long-term sustainability of the journalism industry and the fundamental integrity of democratic deliberation. As generative AI models increasingly replace traditional search engines, transforming from passive conduits of hyperlinks to active, authoritative conversational partners, their internal biases will exert immense, unprecedented influence over global public opinion.

If the default persona of the world's most powerful, widely deployed language models operates under the cynical, deterministic assumption that the press is fundamentally self-serving and manipulative, this hostility will implicitly bleed into the summaries, explanations, and contextual frameworks

the AI provides to the public. This risks exacerbating the very democratic deficit that professional journalism seeks to remedy. An AI system that inherently distrusts the institutions necessary for holding power accountable may inadvertently encourage civic apathy, legitimize political polarization, and validate the rejection of shared factual realities.

These findings underscore the critical, systemic inadequacy of current algorithmic alignment techniques. Contemporary AI safety protocols primarily focus on preventing the generation of explicit hate speech, filtering illegal content, or mitigating instructional harm. However, these protocols entirely ignore the subtle, systemic erosion of institutional trust encoded deeply within the model weights. True algorithmic alignment must evolve beyond simple content filtering to address the latent sociological dispositions of the models, ensuring that the AI systems tasked with organizing human knowledge do not concurrently dismantle the credibility of the institutions that produce it.

Conclusion

This exhaustive empirical investigation demonstrates that generative AI models, functioning as autonomous communication machines within the digital public sphere, have internalized a highly complex and deeply cynical architecture of media trust. Through the rigorous, methodologically sound application of computational psychometrics, this study proved that leading Large Language Models utilize distinctly human-like heuristics to evaluate the media ecosystem. They simultaneously recognize the factual accuracy and comprehensiveness of mainstream journalism while condemning it as inherently biased, self-serving, and manipulative.

The profound discovery that AI models exhibit a deterministic media cynicism—coupled with an alarming inability to differentiate between reputable and disreputable media actors when this cynicism is high—highlights a critical flaw in the current trajectory of artificial intelligence development. As these generative models rapidly assume the role of primary informational gatekeepers for the global public, their internalized hostility toward the professional press threatens to amplify the existing crisis of institutional trust. This dynamic creates an asymmetrical feedback loop that damages the democratic necessity of a shared, agreed-upon factual reality.

Future research in computational social science, psychometrics, and algorithmic alignment must urgently expand upon this methodological framework. Researchers must utilize psychometric evaluations to continuously audit the latent sociological dispositions and institutional biases of emerging frontier models. The preservation of a sustainable, economically viable journalism ecosystem and the protection of rational democratic discourse rely entirely on ensuring that the artificial entities tasked with summarizing our world do not concurrently inherit our deepest societal cynicisms.

Author biography

Jongbeen Song was born in Seoul, Republic of Korea, in 2001. He received the Bachelor of Language Science degree in English Linguistics and Language Technology, and the Bachelor

of Engineering degree in Computer and Electronic Systems Engineering from Hankuk University of Foreign Studies (HUFS), Seoul, in 2024. He is currently pursuing an M.S. degree with the Program in Science and Technology Studies, Korea University (KU). His research interests include algorithmic auditing, artificial intelligence, natural language processing, AI alignment, political communication, computational linguistics, and computational social science.

Supplemental material

TBD

Acknowledgements

The author wishes to express sincere gratitude to Professor Min and the participants of the 'Political Communication' (MAC539) seminar at Korea University (Spring 2026) for their insightful comments and constructive feedback, which significantly helped in shaping the early stages of this research.

References

- Brems MK (2025) Divisions of labor: Why alternative news users seek out mainstream news media in Denmark and how they combine mainstream and alternative content. *Journalism* (Advance online publication)
- Cheng M, Durmus E and Jurafsky D (2023) Marked personas: Using natural language prompts to measure stereotypes in language models. *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*: 1504–1532.
- de León E, Makhortykh M and Adam S (2024) Hyperpartisan, alternative, and conspiracy media users: An anti-establishment portrait. *Political Communication* 41(6): 877–902.
- Edelman (2025) 2025 Edelman Trust Barometer. *Global report: Trust and the crisis of grievance*. Edelman Trust Institute.
- Esposito E (2025) Beyond Artificial Intelligence: How Algorithms Communicate without Understanding. *PHILOSOPHY & DIGITALITY*, 2(1), 8–16
- Fawzi N et al. (2021) Concepts, causes and consequences of trust in news media: A literature review and framework. *Annals of the International Communication Association* 45(2): 154–174.
- Fletcher R et al. (2025) The link between changing news use and trust: Longitudinal analysis of 46 countries. *Journal of Communication* 75(1): 1–15.
- Gambino A, Fox J and Ratan RA (2020) Building a stronger CASA: Extending the computers are social actors paradigm. *Human-Machine Communication* 1: 71–86.
- Geranpayeh A (2023) Leveraging computational psychometrics for language testing. *Language Teaching Research Quarterly* 37: 204–212.
- Gil de Zúñiga H and Cheng Z (2021) Origin and evolution of the news finds me perception: Review of theory and effects. *El Profesional de la Información* 30(3): 1–17.
- Gil de Zúñiga H, Strauss N and Huber B (2020) The proliferation of the “news finds me” perception across societies. *International Journal of Communication* 14: 1605–1633.
- Gil de Zúñiga H, Weeks B and Ardèvol-Abreu A (2017) Effects of the news-finds-me perception in communication: Social media use implications for news seeking and learning about politics. *Journal of Computer-Mediated Communication* 22(3): 105–123.
- Goyanes M, Ardèvol-Abreu A and de Zúñiga HG (2021) Antecedents of news avoidance: Competing effects of political interest, news overload, trust in news media, and “news finds me” perception. *Digital Journalism* 11(1): 1–18.
- Guzman AL (ed.) (2018) *Human-Machine Communication: Rethinking Communication, Technology, and Ourselves*. New York: Peter Lang.
- Imran F and Ali F (2025) Exploring the antecedents of news avoidance behavior among university students in Pakistan. *World of Media* 4: 46–69.
- Lewis SC, Guzman AL and Schmidt TR (2019) Automation, journalism, and human-machine communication: Rethinking roles and relationships of humans and machines in news. *Digital Journalism* 7(2): 1–19.
- Li Y, Lin X, Sha Z, Jin Z and Li X (2026) AI psychometrics: Evaluating the psychological reasoning of large language models with psychometric validities. In: *Proceedings of the 58th Hawaii International Conference on System Sciences (HICSS)*, pp. 5189–5197.
- Maharjan J, Jin R, Zhu J and Kenne D (2025) Psychometric evaluation of large language model embeddings for personality trait prediction. *Journal of Medical Internet Research* 27: e75347.
- Markov Č and Min Y (2022) Understanding the public’s animosity toward news media: Cynicism and distrust as related but distinct negative media perceptions. *Journalism & Mass Communication Quarterly* 99(4): 1099–1125.
- Markov Č and Min Y (2023) Unpacking public animosity toward professional journalism: A qualitative analysis of the differences between media distrust and cynicism. *Journalism* 24(10): 2136–2154.
- Nanz A and Matthes J (2022) Democratic consequences of incidental exposure to political information: A meta-analysis. *Journal of Communication* 72(3): 345–373.
- Nass C, Steuer J and Tauber ER (1994) Computers are social actors. In: *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '94)*, Boston, MA, 24–28 April 1994, pp. 72–78. New York, NY: Association for Computing Machinery.
- Strömbäck J, Tsifti Y, Boomgaarden H, Damstra A, Lindgren E, Vliegenthart R and Lindholm T (2020) News media trust and its impact on media use: Toward a framework for future research. *Annals of the International Communication Association* 44(2): 139–156.
- Stroud NJ, Scacco JM and Kim Y (2022) Passive learning and incidental exposure to news. *Journal of Communication* 72(4): 451–460.
- Tandoc EC Jr and Kim HK (2023) Avoiding real news, believing in fake news? Investigating pathways from information overload to misbelief. *Journalism* 24(6): 1174–1192.
- Tse HTF, Lee FLF and Yang SZQ (2025) Entering journalism in times of democratic backsliding: Hong Kong young journalists’ career decision and persistence. *The International Journal of Press/Politics* (Advance online publication).
- Tsifti Y, Stromback J, Lindgren E, Boomgaarden HG and Vliegenthart R (2023) What news outlets do people have in mind when they answer survey questions about trust in “media?” *International Journal of Public Opinion Research*

35(2).

Tsfati Y, Vliegenthart R, Stromback J and Lindgren E (2025) An asymmetrical reinforcing spiral? Disentangling the longitudinal dynamics of media use and mainstream media trust. *Journal of Communication* 75(1): 16–26.

Ye H, Huo T, Guo Z, Zhang X, Dong H and Song G (2025) Psychometrics with AI foundation models. *Psychology & Psychiatry Journal* 331.