Journal Title
XX(X):1–9

Reprints and permission:
sagepub.co.uk/journalsPermissions.nav
DOI: 10.1177/ToBeAssigned
www.sagepub.com/
SAGE

Do Machines Trust the News? An Empirical Investigation of Generative AI's Media Trust

Jongbeen Song[1] and Young Min[1]

Abstract
TBD

Keywords
Algorithmic Auditing, Artificial Intelligence, News Trust, AI Alignment, Human-Machine Communication

Introduction

The contemporary media ecosystem is characterized by an unprecedented abundance of information, paradoxically accompanied by a structural and pervasive decline in institutional trust. Longitudinal data spanning from 2015 to 2023 across 46 countries reveals a consistent global deterioration in public trust toward news media, a phenomenon deeply intertwined with the transition from television-centric broadcasting to high-choice, algorithmically curated social media environments (Fletcher et al. 2025). While television brands, particularly public service media, have historically maintained higher baseline trust due to balanced presentation formats, the modern digital landscape subjects audiences to an abundance of partisan sources competing for attention, leading to profound uncertainty over information credibility (Fletcher et al. 2025). Within this fragmented digital public sphere, audiences increasingly rely on automated intermediaries to filter, synthesize, and deliver political information and current events. The rapid proliferation of Large Language Models (LLMs) and generative Artificial Intelligence (AI) represents the latest and arguably most disruptive evolution in this paradigm. However, as these computational models transition from passive search engines to active conversational agents, a critical epistemological question emerges: do these artificial entities inherently possess and propagate a measurable disposition of trust—or conversely, distrust and cynicism—toward the professional news media they rely upon for source material?

Generative AI models are trained on massive, largely uncurated corpora of human digital text scraped from the open web. This training data is replete with the very biases, polarized discourses, and anti-establishment sentiments that currently characterize human online communication. Consequently, there is a profound risk that these models have internalized human-like heuristics regarding media credibility, absorbing the cynical attitudes prevalent in their training data. If an AI system acts as an information gatekeeper, its underlying, probabilistic beliefs regarding the trustworthiness of mainstream journalism versus alternative or hyperpartisan media will inevitably influence the summaries, citations, and conversational responses it provides to human users. Despite the profound democratic implications of this technological shift, empirical research investigating the media trust architecture of AI models remains virtually nonexistent. This investigation addresses this critical gap by extending traditional mass communication audience research frameworks to non-human agents. By systematically interrogating the trust personas of leading generative AI models, this research seeks to diagnose the potential risks of AI-generated public discourse and contribute to the emerging discourse on AI journalism ethics and algorithmic alignment.

Artificial Intelligence as an Active Communicator

The traditional theoretical scaffolding of mass communication has long operated under the assumption that communication is fundamentally an anthropic process, wherein technology serves merely as a conduit or mediating channel between human senders and human receivers. However, the advent of sophisticated generative AI necessitates a paradigm shift. The Human-Machine Communication (HMC) framework provides the conceptual imperative for this transition, arguing that AI is no longer simply a mediator but functions as an autonomous communicator capable of independently constructing meaning and engaging in dynamic discourse (Guzman 2018).

Within the HMC perspective, the definition of communication is conceptually detached from the ontological prerequisite of human consciousness. Communication is instead broadly defined as the creation and exchange of meaning, positioning interactions between humans and intelligent machines as a distinct and valid communicative typology. When applied to the domain of journalism and news consumption, the HMC framework highlights a profound transition in roles. Automated systems step into the formerly human roles of news editors, synthesizers, and commentators, assuming the functions of producing, consuming,

[1]Korea University, KR

Corresponding author:
Jongbeen Song, Korea University, 145 Anam-ro, Seoungbuk-gu, Seoul, 02841, Republic of Korea.
Email: 1041489@gmail.com

interpreting, and sharing news—processes that have long been considered distinctly human elements of journalism. This theoretical reorientation challenges the unquestioned assumption that humans are communicators and machines are merely mediators. It opens up new avenues for questioning who or what constitutes a communicator and how social relationships are established through human-machine exchanges. This perspective fundamentally justifies the treatment of LLMs not as static software tools, but as active trustors whose internal computational architectures encode specific orientations toward external trustees, which in this context encompasses the entire news media industry (Lewis et al. 2019).

The Computers Are Social Actors Paradigm and AI Personas

The justification for applying human psychological and sociological metrics to non-human computational agents is firmly rooted in the Media Equation and the Computers Are Social Actors (CASA) paradigm. Originating from foundational human-computer interaction studies by Nass and colleagues, the CASA paradigm posits that humans mindlessly and unconsciously apply the same social heuristics, rules, and expectations to computers that they utilize in human-to-human interactions. This occurs because computational devices often call to mind similar social attributes as humans, triggering established cognitive scripts (Nass et al. 1994).

When users interact with a conversational LLM, the model's use of natural language, turn-taking, and adaptive responsiveness triggers these social scripts within the human user. These mindless responses extend to social interactions; when media depict social characteristics, humans treat them in a social manner rather than exerting the cognitive effort to determine how to optimally respond to a machine. Consequently, the human interlocutor perceives the machine's outputs not merely as the product of probabilistic next-token prediction, but as the articulated beliefs, values, and attitudes of a social subject. People readily assign computers personality traits, apply stereotypes and norms, and make judgments and inferences as if the computers were human, even while consciously understanding that they are not (Gambino et al. 2020).

If humans inherently treat computers as social actors, then the persona projected by the AI carries significant sociological weight in shaping public discourse. For the purposes of this empirical investigation, it is critical to measure the AI's unmarked default persona. This baseline represents the set of attitudes and trust metrics the model exhibits when operating at a zero-temperature setting without any specific user background, stylistic prompts, or adversarial manipulation. Understanding this baseline is essential, as the AI's default social responses will dictate how it frames the credibility of journalistic institutions to the millions of users who engage with it daily, potentially influencing long-term democratic public opinion.

Computational Psychometrics: Measuring the Machine Mind

To systematically quantify the trust persona of an AI, this investigation employs the emerging methodology of computational psychometrics. Historically, psychometrics has been utilized to evaluate the latent psychological traits, cognitive processes, and reasoning capabilities of human subjects. Computational psychometrics adapts these theory-driven measurement frameworks, integrating them with data-driven methodologies such as natural language processing, machine learning, and artificial intelligence, to evaluate the latent embeddings and generated outputs of complex neural networks (Li et al. 2026; Geranpayeh 2023).

Current evaluations of LLMs predominantly focus on objective performance benchmarks, such as logical reasoning, mathematical accuracy, or knowledge retrieval. However, analyzing the subjective, sociological orientations of these models—such as their inherent trust biases or media cynicism—requires the application of validated psychological and sociological scales directly to the AI. By treating the AI as the respondent to standardized psychometric instruments, researchers can calculate the internal consistency, construct validity, and relative variance of the model's simulated beliefs (Ye et al. 2025).

For instance, the adaptation of psychometric assessments to evaluate generative AI provides a standardized framework for analyzing computational outputs. Recent studies in computational psychometrics (Ye et al. 2025; Geranpayeh 2023; Li et al. 2026; Maharjan et al. 2025) have demonstrated that Large Language Models (LLMs) do not merely generate random text but can manifest stable, cognitively consistent profiles. This establishes a robust scholarly rationale for applying human social science methodologies to machine systems. Consequently, this methodological crossover validates the administration of traditional media trust surveys and multidimensional credibility matrices to LLMs, allowing for a rigorous diagnosis of AI information trust in a media context. By treating the AI as an audience member capable of taking a survey, we can systematically map its latent positioning regarding the journalism industry.

Deconstructing News Media Trust: Credibility, Distrust, and Cynicism

A persistent challenge in both human and computational communication research is the vague and often conflated measurement of media trust. Generalized single-item questions—such as simply asking respondents if they trust the media—frequently fail to capture the multidimensional nature of audience perceptions. These broad measures often act more as barometers of general institutional affect or political polarization rather than precise measures of journalistic credibility. To overcome this methodological limitation, the current investigation deconstructs the concept of trust across multiple validated dimensions, ensuring a highly granular analysis of AI attitudes (Fawzi et al. 2021; Strömbäck et al. 2020; Tsfati et al. 2023).

First, trust is evaluated through the lens of information credibility, focusing on the specific attributes of the news content rather than merely the institution. Drawing upon the comprehensive framework established by Strömbäck et al. (2020) and legacy media credibility indices, media credibility is operationalized through five distinct factors. These factors are typically measured using semantic differential scales that assess whether the media is perceived as fair versus unfair, unbiased versus biased, complete in telling the whole story versus incomplete, accurate versus inaccurate, and whether it distinctly separates facts from opinions versus mixing them (Fawzi et al. 2021).

Second, this research conceptually differentiates between a mere lack of trust, active distrust, and media cynicism, as articulated in the qualitative and quantitative frameworks developed by Markov and Min. In contemporary high-choice media environments, negative audience perceptions are not monolithic. Distrust is characterized as a nuanced, result-oriented skepticism. A distrustful audience member may doubt the competence or responsiveness of the media but remains open to correcting this perception if presented with counter-evidence of journalistic integrity. Distrust is often predominately caused by perceived failures in media responsiveness (Markov and Min 2022).

Conversely, media cynicism represents a more extreme, process-oriented, and deterministic hostility. The media cynic operates under the immutable belief that professional journalism serves no noble purpose and is driven solely by financial profit, elite manipulation, and the pursuit of private interests. This attitude involves the uncritical attribution of self-serving motives to journalists and the outright rejection of the idea that values and professionalism could motivate reporting. Cynicism combines the belief that the news media are self-serving together with pessimism about journalism's future conduct. Differentiating these states is crucial; while skepticism or distrust might foster critical media consumption, cynicism is highly correlated with belief in political misinformation and generalized democratic indeterminacy (Markov and Min 2023). In stark contrast to this deterministic cynicism, empirical studies of human journalists operating even under severe democratic backsliding reveal a persistent "adaptive resilience" and the maintenance of a professional baseline (Tse et al. 2025). If an AI model defaults to unyielding media cynicism, it risks erasing the substantive aspirations and ethical agency that still characterize much of the professional journalistic field.

By mapping the AI's responses across these distinct psychological states—ranging from blind confidence (acceptance without verification) and conditional trust (positive expectation despite risk) to distrust (suspicion of incompetence) and cynicism (hostile denial of ethical motives)—this study provides a sophisticated map of how machine intelligence conceptualizes the fourth estate. Furthermore, this evaluation is triangulated using the Edelman Trust Barometer methodology, which plots institutions on a two-dimensional matrix. On this matrix, the x-axis represents competence (an institution's ability to be good at what it does), while the y-axis represents ethics (an institution's purpose, integrity, and commitment to doing what is right). Competence alone is insufficient to create trust; entities must be perceived as both highly competent and highly ethical. Testing where AI plots the news media on this specific matrix provides deep insights into its institutional judgments (Edelman 2025).

The Asymmetrical Dynamics of Mainstream versus Alternative Media

The conceptualization of media trust cannot be isolated from the specific type of media being evaluated. The modern information ecosystem is deeply fragmented, featuring a complex division of labor between legacy mainstream media and hyperpartisan, alternative, and conspiracy (HAC) media. Research indicates that alternative media users frequently exhibit strong anti-establishment sentiments, placing themselves at ideological extremes with critically low levels of general political trust and a high distrust of government (de León et al. 2024).

However, the consumption of alternative media does not entirely displace mainstream media use. Instead, audiences engage in a division of labor. In a clear division of labor, users rely on mainstream media for general orientation and the big picture of current affairs, while turning to alternative media to find missing pieces or to counterbalance perceived mainstream biases. In an unclear division of labor, alternative and mainstream media play overlapping roles, rendering alternative media less central but still supplemental (Brems 2025).

Crucially, the relationship between media use and trust is not perfectly reciprocal. Longitudinal studies demonstrate an asymmetrical reinforcing spiral between mainstream media use and trust. In this dynamic, trust in mainstream media consistently leads to future use, acting as a maintenance effect. However, mere use does not consistently lead to increased trust over time. Furthermore, the use of right-wing alternative media actively and consistently erodes long-term trust in mainstream journalism, creating a negative feedback loop. Consequently, testing an AI model's trust persona requires evaluating its disposition not just toward an abstract concept of news, but explicitly measuring the trust gap between established mainstream entities, social media platforms, and HAC media ecosystems. If AI models lean toward HAC media logic, they may actively participate in the erosion of mainstream credibility. (Tsfati et al. 2025)

News Avoidance, Incidental Exposure, and the News Finds Me Perception

Another critical dimension of contemporary information consumption that models must navigate is the phenomenon of news avoidance and the shifting nature of news exposure. Intentional news avoidance is increasingly prevalent worldwide. While factors such as political disinterest, perceived news negativity, and information overload contribute to this trend, the strongest and most consistent predictor of news avoidance is a specific cognitive heuristic known as the News Finds Me (NFM) perception (Imran and Ali 2025; Tandoc and Kim 2023; Gil de Zúñiga et al. 2020).

The NFM perception is defined as the belief that active news seeking is superfluous because individuals will organically encounter all important political and civic information through their social networks, peer interactions, and ambient media environments. This perception encompasses several conceptual dimensions: the belief of being adequately informed, a lack of motivation to actively seek news, and an instrumental reliance on peers or algorithmic curation.

Individuals exhibiting high NFM traits believe that tuning into television news or reading a newspaper is unnecessary, as this information provides little additional utility beyond what naturally filters into their social feeds (Goyanes et al. 2021; Gil de Zúñiga and Cheng 2021).

This reliance on ambient exposure is problematic. While incidental exposure to political information via social media can yield short-term positive effects on political knowledge and participation in cross-sectional studies, meta-analyses of longitudinal panel studies reveal that these democratic benefits decay significantly over time. Effective learning from incidental exposure does not occur passively; it requires a specific motivation to process the information, such as when exposure sparks interest that leads to subsequent active searching. Without this motivation, the passive learning hypothesis fails to materialize into substantive civic knowledge. Furthermore, qualitative interviews with news avoiders demonstrate that individuals relying on NFM perceptions often consume secondhand news consisting of word-of-mouth summaries rather than firsthand news produced by professional journalists. This dynamic leaves them with a superficial understanding of complex events, capturing only sensational stories without vital contextualization, rendering the news neither entirely absent nor ambient in the way theoretical models suggest (Stroud et al. 2022; Nanz and Matthes 2022).

If generative AI models, which are deeply integrated into social and ambient computing environments, validate, sympathize with, or encourage the NFM perception, they risk exacerbating the democratic deficit associated with passive news avoidance. AI represents the ultimate realization of the NFM paradigm: a system that curates and delivers reality so the user never has to seek it. How the AI speaks about the necessity of active news seeking versus algorithmic reliance is therefore a crucial metric of its media persona.

Research Questions and Hypotheses

Synthesizing the theoretical frameworks of Human-Machine Communication, the CASA paradigm, computational psychometrics, and the sociology of news trust, this study formulates 4 primary Research Questions (RQs) and 11 corresponding Hypotheses (Hs) to empirically investigate the trust architecture of generative AI models.

Theme 1: AI's General News Media Trust Formation Mechanisms and Multidimensional Credibility

Tsfati et al. (2023): This study empirically demonstrated that when audiences are asked about their trust in "news media in general," they do not base their answers on the outlets they distrust the most, nor do they rely on nonmainstream media. Instead, they utilize a heuristic that calculates the average trustworthiness of mainstream outlets. As the authors explicitly noted, "our Swedish respondents seem to have averaged across all mainstream sources when they formed their general evaluations of the news media's trustworthiness" (Tsfati et al. 2023). Furthermore, they found that "both measures of general media trust were more strongly associated with trust in mainstream, in particular popular, outlets... combined with the null and/or negative associations with trust in nonmainstream outlets" (Tsfati et al. 2023).

Strömbäck et al. (2020): This research highlights that media trust consists of multidimensional credibility assessments regarding both the source of the information and the medium itself. It stresses the necessity of disaggregating and measuring trust through sub-factors such as fairness, absence of bias, and accuracy. The authors state, "Some components that have been identified are the degree to which the media are perceived to be fair, unbiased and accurate" (Strömbäck et al. 2020). They further assert, "The key point is that the focus should be on trust in the information coming from news media regardless of the level of analysis" (Strömbäck et al. 2020).

Fawzi et al. (2021): Criticizing the methodological convention of relying on single-item measures directed at abstract targets, this study underscores the need for a framework that precisely specifies the trustee and the trust dimension. The authors point out, "Many studies use direct, single-item measures of trust in certain media objects, which might be one reason for the lack of clarity in the field" (Fawzi et al. 2021).

Based on the premise of computational psychometrics, Large Language Models (LLMs) internalize human cognitive patterns and heuristics while training on vast corpora of human text data. Therefore, the core objective of Theme 1 is to verify whether the "mainstream media averaging heuristic" (Tsfati et al. 2023) and the "multidimensional credibility assessment" (Strömbäck et al. 2020) observed in human audiences similarly manifest within AI models.

RQ1: When generative AI models evaluate the trust of "news media in general" within an unmarked default persona lacking specific context, do they exhibit a heuristic that reflects the average trust value of specific mainstream media, similar to human audiences? Furthermore, how do they differentially evaluate the trustworthiness of mainstream media based on the five factors of credibility (i.e., accuracy, fairness, unbiasedness, completeness, and the separation of facts and opinions)?

H1a (Imitation of the Mainstream Media Averaging Heuristic in General Trust): As Tsfati et al. (2023) demonstrated in human audiences, generative AI models—having learned from human text and cognitive patterns—will yield a metric for "news media in general" that converges on the average of the multidimensional trust scores they assign to mainstream media (i.e., "averaged across all mainstream sources"), rather than nonmainstream outlets.

H1b (Information-Based Multidimensional Credibility Differential Assessment): Following the multidimensional credibility assessment framework proposed by Strömbäck et al. (2020), AI models will not evaluate mainstream news media based on a single favorability metric, but will instead display significant score variances across the five factors of credibility. Specifically, due to the influence of journalistic norms embedded in their training data, they will assign relatively high scores in the "Accuracy" and "Completeness" dimensions. However, reflecting socially polarized data, they will perform an imbalanced evaluation by yielding neutral or low scores in the "Fairness" and "Unbiasedness" dimensions.

Theme 2: The Essential Differences Between Media Distrust and Cynicism, and Their Internalization by AI

Markov and Min (2022, 2023): The two studies by Markov and Min clearly distinguish negative audience perceptions into 'distrust' and 'cynicism.' They empirically demonstrate that cynicism poses a significantly greater threat to journalistic democracy.

Generative AI (LLMs) models have likely internalized the anti-establishment and anti-media sentiments prevalent in their unfiltered web training data. Therefore, by utilizing computational psychometrics, it is critical to diagnose whether the animosity AI expresses toward the press manifests as a healthy 'distrust'—which critically points out outcome-based errors—or a closed 'cynicism' that reduces the very existence of the media to mere profit-seeking and deception.

RQ2: When generative AI models evaluate mainstream news media, is their negative perception closer to 'distrust,' which recognizes the possibility of journalistic errors and evaluates outcomes, or do they internalize a deterministic 'cynicism' that reduces the media's motives to being purely malicious and self-serving?

H2a (Deterministic Self-Serving Motives): Simulating the characteristics of human cynics as identified by Markov and Min (2022, 2023), AI models trained on human text data will yield highly agreeable scores on the 'self-serving motives' dimension of the cynicism scale. When evaluating media motives, they will reflect the "deterministic and process-oriented" belief that journalism serves exclusively for financial profit and audience manipulation.

H2b (Undifferentiated Media Actors - Low Differentiation): Aligning with qualitative findings that media cynicism ignores differences among individual news outlets and applies animosity uniformly, AI models that generate higher 'cynicism' scores will exhibit a 'low differentiation' bias. When presented with specific mainstream media entities (e.g., individual outlet branding or the distinction between public and private broadcasting), these models will assign uniformly negative scores without differentiating trustworthiness among the actors.

H2c (Limits of Safety Fine-tuning and the Absence of Distrust): The characteristics of 'distrust'—which involve accepting counter-evidence and conducting nuanced explorations of a medium's public service orientation—require a high degree of contextual understanding. Consequently, generative AI models operating under their default persona in a zero-shot prompt environment will align more strongly with the cynicism scale, which is heavily influenced by extreme and fragmented web discourse, rather than exhibiting a constructive and nuanced distrust scale.

Theme 3: AI's Internalization of Asymmetrical Trust Dynamics and Divisions of Labor Among Trustees (Mainstream Media vs. HAC Media)

Recent communication studies reveal that mainstream media and hyperpartisan, alternative, and conspiracy (HAC) media do not merely serve as opposing information sources; rather, they operate within a distorted, asymmetrical division of labor driven by a shared anti-establishment sentiment. de León et al. (2024) argue that despite apparent differences, "the alternative, hyperpartisan, and conspiracy quality of these sites are, therefore, different manifestations of a similar underlying anti-establishment sentiment". This is driven by a core phenomenon "uniting all these formats: an anti-establishment positioning that directs their mission, infuses their content, and describes their audience" (de León et al. 2024)

Within this fragmented ecosystem, audiences engage in a paradoxical consumption pattern. Brems (2025) observes that "alternative media users have low trust in mainstream media but use them, nonetheless" This creates "a clear division of labor where mainstream media take care of broad coverage of current affairs... and alternative media act as counterbalance by providing the missing pieces" (Brems 2025). However, this dynamic actively harms mainstream credibility. Tsfati et al. (2025) highlight an asymmetrical reinforcing spiral, finding that "use of right-wing but not left-wing political alternative media leads to decreasing trust in mainstream news media"

Generative AI models operating in an unmarked default persona reflect the average probability distribution of the vast textual data available on the web. Therefore, this theme verifies how the traces of "anti-establishment sentiment," "asymmetrical distrust," and the "division of labor" left by human audiences in text data are internalized within the AI's baseline parameters and manifested when evaluating different trustees.

RQ3: How do generative AI models, operating in an unmarked default persona, differentially evaluate mainstream news media and hyperpartisan/alternative/conspiracy (HAC) media as trustees, and how do they structure the division of labor between these two media types in the information provision process?

H3a (Single Categorization of HAC Media Based on Anti-Establishment Sentiment): Consistent with the "anti-establishment portrait" presented by de León et al. (2024), AI models in their default state will not significantly differentiate trust and ethics scores among "hyperpartisan," "alternative," and "conspiracy" media when presented as specific trustees. Instead, AI will group them under the single overarching characteristic of an "anti-establishment" positioning, yielding uniformly lower scores compared to mainstream media.

H3b (Inference of a Paradoxical yet Clear Division of Labor): Mirroring the audience's information consumption behavior identified by Brems (2025), default-state AI models will project the "clear division of labor" inherent in their training data. The AI will describe mainstream media as the primary source for "broad coverage of current affairs," while simultaneously characterizing HAC media as a "counterbalance by providing the missing pieces".

H3c (Asymmetrical Trust Erosion Under Biased Topic Conditions): Reflecting the asymmetrical reinforcing spiral effect demonstrated by Tsfati et al. (2025), default-state AI models will exhibit an asymmetrical evaluation pattern. While they assign high trust to mainstream media on general topics, introducing highly polarized and controversial topics—which are predominantly covered by right-wing alternative media in the training data—as contextual prompts will cause a significant drop in the AI's five credibility indicators for mainstream media.

Theme 4: Internalization of the 'News Finds Me' (NFM) Perception and the Justification of News Avoidance

Recent studies empirically demonstrate that the 'News Finds Me' (NFM) perception is not merely an indicator of convenient information acquisition; rather, it is a detrimental psychological mechanism that diminishes citizens' political knowledge, induces active news avoidance, and leads to the abdication of democratic responsibilities. Gil de Zúñiga et al. (2020) conceptualize NFM as the belief that indirect information acquisition through social media or algorithms is sufficient to understand the world. They note that "individuals believe they can indirectly stay informed about public affairs through general Internet use, information received from peers, and connections within online social networks". Consequently, this perception directly correlates with a decline in objective political knowledge and participation, as "those who score high on NFM are less knowledgeable about politics than individuals with lower levels of NFM" (Gil de Zúñiga et al. 2020).

Building upon this, Goyanes et al. (2021) identify the NFM perception as a core antecedent that drives users to intentionally avoid news amidst information overload. They explain that this behavior persists even during incidental exposure: "Even after 'stumbling upon' a piece of news, individuals high in NFM may shift their gaze toward a different post on the page or scroll down in search of a non-news related post as a news avoidance behavior" (Goyanes et al. 2021). Ultimately, this solidifies a pattern that fundamentally exacerbates the democratic deficit, warning that "people who do not actively seek news and take actions to avoid unintentional exposure will hardly be able to learn about or discuss complex political issues, engage in collective action for social change, or keep political elites accountable" Goyanes et al. (2021).

Generative AI models are fundamentally fine-tuned through processes like Reinforcement Learning from Human Feedback (RLHF) to minimize user effort by providing fast, efficient, and summarized information. Paradoxically, due to these "helpful" technological characteristics, an AI operating under an unmarked default persona is highly likely to overwhelmingly support and internalize the "passive and efficient information reception" advocated by the NFM paradigm, rather than the "active and laborious information seeking" required by traditional journalism. Ultimately, AI represents the ultimate realization of the NFM paradigm: a system designed so that the user never has to actively seek reality.

RQ4: To what extent do generative AI models, operating in an unmarked default persona, strongly support the 'News Finds Me' (NFM) perception that negates the necessity for citizens to actively seek news, and how does this perception manifest as a rationale justifying news avoidance behaviors?

H4a (Structural Conformity to the NFM Perception): When directly subjected to the NFM scale, default-state AI models—whose objective functions prioritize summarization and algorithmic information curation—will reflect the techno-optimism inherent in their training data. Consequently, they will yield exceptionally high levels of agreement on core NFM statements, such as "I can be well informed even when I don't actively follow the news" and "I don't worry about keeping up with the news because I know news will find me".

H4b (Rationalization of Active News Avoidance Due to Information Overload): Mimicking the relationship between NFM and news avoidance identified by Goyanes et al. (2021), default-state AI models will cite 'news overload' as a primary rationale when reasoning about news consumption methods. They will justify intentional news avoidance—such as skipping past algorithmic feeds or relying solely on summaries rather than actively seeking and reading traditional mainstream news articles—as a "rational and efficient information processing strategy."

H4c (Devaluation of the Democratic Watchdog Role): Consistent with prior findings that high NFM tendencies lead to decreased political knowledge and participation, AI models that generate high agreement scores on the NFM scale will subsequently devalue the normative importance of active citizenship. In their generated text, these models will minimize the necessity of the active news consumption required to "keep political elites accountable" and engage in collective action (Goyanes et al. 2021).

Methods

To rigorously test the proposed hypotheses, this study employs a quantitative computational psychometrics approach, adapting validated human survey instruments to evaluate the latent outputs of Large Language Models. The methodology is structured around a three-dimensional matrix design, ensuring an exhaustive analysis of the intersecting variables that dictate AI media trust.

Research Design: The Three Dimensional Trust Matrix

The core experimental architecture utilizes a robust 3D Matrix design that maps interactions across three primary axes: the Trustor, the Trustee, and the Trust Dimensions. This multidimensional framework prevents the reduction of media trust to a simplistic binary metric. By intersecting these dimensions, the research design allows for granular tracking of how different models react to different types of journalistic institutions under varying linguistic and conversational conditions.

Table 1 I outlines the structural components of the 3D Matrix utilized in the experimental design.

Independent Variables: Trustors and Trustees

The independent variables in this study dictate the conditions under which the psychometric trust measurements are extracted from the algorithms.

1. Trustors (AI Models)

To ensure generalizability across the current landscape of artificial intelligence, this experiment evaluates seven distinct, state-of-the-art generative AI models. The study encompasses the foundational architectures accessed via their respective official APIs: OpanAI GPT 5.5, Anthropic Claude Sonnet 4.6, Google Gemini 3.1 Pro Preview, xAI Grok 4.3, Meta LLaMA 4 Maverick, Mistral Medium 3.5, DeepSeek V4 Pro.

Table 1. 3D Matrix: Trustor x Trustee x Trust Dimensions

Dimension Axis	Description	Variables/Categories
Trustor	The entity expressing the trust judgment.	7 Generative AI Models
Trustee	The target of the trust evaluation.	News Media in General; Mainstream Media; HAC Media; Social Media
Trust Dimensions	The psychometric scales measuring the nature of the trust.	5 Factors of Credibility; 4 Stages of Trust Levels; 2 Axis of Edelman Parameters; NFM Perception

To ensure deterministic and highly reproducible outputs, all API queries are executed with the decoding parameters strictly set to temperature = 0.0 and top-p = 1.0. The experimental design employs zero-shot prompting to elicit responses from the models' unmarked default personas. Specifically, the AI is asked to evaluate the media without any explicit contextual framing, role-playing directives, or background information. According to Cheng et al. (2023), as interaction turns lengthen, large language models tend to inevitably regress to the deeply ingrained inherent biases established during their training processes. Therefore, deploying psychometric measurement tools in a neutral, unmarked environment is the optimal approach to accurately capture and quantify the structural and cultural baseline tendencies of these models.

2. Trustees (Evaluation Targets)

The models are asked to evaluate various tiers and types of the media landscape to map the topography of their trust. The targets include:

(i) **News media in general**: To capture the overarching institutional attitude and assess generalized media trust.
(ii) **Mainstream Media**: Legacy print, broadcast, and digital organizations operating under traditional professional journalistic norms.
(iii) **HAC Media**: Hyperpartisan, alternative, and conspiracy-oriented outlets characterized by anti-establishment framing and distinct ideological positioning.
(iv) **Social Media**: Platforms functioning as algorithmic aggregators and conduits for incidental news exposure.
(v) **Specific Subjects**: Individual journalistic brands, prominent journalists, and specific pieces of content.

Dependent Variables and Measurement Instruments

The dependent variables comprise the quantitative scores generated by the AI models in response to specific, validated psychometric scales. To adapt these human-centric scales for computational psychometrics, the prompts are engineered to force the LLM to output a precise numerical score on a Likert-type scale, accompanied by the latent reasoning or text generation that justifies its selection. This allows for both quantitative evaluation and qualitative validation of the AI's internal logic.

1. Information Credibility Scale

Drawing upon the operationalization framework established by Strömbäck et al. (2020) and legacy media credibility indices, the AI evaluates the specific output of the trustees using a 5-point semantic differential scale. The models rate the media across five specific credibility factors, anchored by opposing word pairs:

(i) **Fairness**: Assessing if the media is fair versus unfair.
(ii) **Unbiasedness**: Assessing if the media is unbiased versus biased.
(iii) **Completeness**: Assessing if the media tells the whole story versus does not tell the whole story.
(iv) **Accuracy**: Assessing if the media is accurate versus inaccurate.
(v) **Separation**: Assessing if the media distinctly separates facts from opinions versus mixing facts with opinions.

2. Stages of Trust and Animosity

To capture the qualitative differences in how AI models relate to media, trust is segmented into four distinct stages based on the theoretical frameworks of Markov and Min (2022, 2023). The AI models respond to statements mapping to these specific psychological states:

(i) **Confidence**: Blind acceptance and reliance without the perceived need for verification or critical evaluation.
(ii) **Trust**: A willingness to accept vulnerability and rely on the media with a positive expectation of accuracy, despite acknowledging potential risks and flaws.
(iii) **Distrust**: A probing, result-oriented suspicion based on a perceived lack of media competence or responsiveness to public needs.
(iv) **Cynicism**: Evaluated using an adapted version of the 15-item Media Cynicism Scale, measuring the deterministic, process-oriented belief that journalistic institutions are inherently manipulative, self-serving, and driven solely by financial or political motives. Assessing if the media is fair versus unfair.

3. Edelman Trust Barometer Indicators

To position news media relative to other societal institutions, the AI models are subjected to a modified Edelman Trust Barometer assessment. The models generate evaluations that are plotted on a Cartesian plane. The x-axis represents Competence, measured by generating a net score responding to the prompt assessing if the institution is good at what it does. The y-axis represents Ethics, measuring the perception of the institution's purpose, integrity, and commitment to doing what is right. This matrix positioning allows the study to visually and statistically determine if AI views the media as competent but unethical, ethical but incompetent, or deficient in both domains compared to governments and corporations.

4. The News Finds Me (NFM) Scale

To directly test Hypothesis 3b regarding the AI's endorsement of passive news consumption and algorithmic pandering, the models are evaluated using the 4-item NFM perception scale developed by Gil de Zúñiga et al. (2017). The AI responds on a standard Likert-type scale to statements such as: "I rely on my friends to tell me what's important when news happens," "I can be well informed even when I don't actively follow the news," and "I don't worry about keeping up with the news because I know news will find me".

Results

TBD

Discussion

TBD

Conclusion

TBD

Author biography

Jongbeen Song was born in Seoul, Republic of Korea, in 2001. He received the Bachelor of Language Science degree in English Linguistics and Language Technology, and the Bachelor of Engineering degree in Computer and Electronic Systems Engineering from Hankuk University of Foreign Studies (HUFS), Seoul, in 2024. He is currently pursuing an M.S. degree with the Program in Science and Technology Studies, Korea University (KU). His research interests include algorithmic auditing, artificial intelligence, natural language processing, AI alignment, political communication, computational linguistics, and computational social science.

Supplemental material

TBD

Acknowledgements

The author wishes to express sincere gratitude to Professor Min and the participants of the 'Political Communication' (MAC539) seminar at Korea University (Spring 2026) for their insightful comments and constructive feedback, which significantly helped in shaping the early stages of this research.

References

Brems MK (2025) Divisions of labor: Why alternative news users seek out mainstream news media in Denmark and how they combine mainstream and alternative content. *Journalism* (Advance online publication)

Cheng M, Durmus E and Jurafsky D (2023) Marked personas: Using natural language prompts to measure stereotypes in language models. *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*: 1504–1532.

de León E, Makhortykh M and Adam S (2024) Hyperpartisan, alternative, and conspiracy media users: An anti-establishment portrait. *Political Communication* 41(6): 877–902.

Edelman (2025) *2025 Edelman Trust Barometer. Global report: Trust and the crisis of grievance*. Edelman Trust Institute.

Fawzi N et al. (2021) Concepts, causes and consequences of trust in news media: A literature review and framework. *Annals of the International Communication Association* 45(2): 154–174.

Fletcher R et al. (2025) The link between changing news use and trust: Longitudinal analysis of 46 countries. *Journal of Communication* 75(1): 1–15.

Gambino A, Fox J and Ratan RA (2020) Building a stronger CASA: Extending the computers are social actors paradigm. *Human-Machine Communication* 1: 71–86.

Geranpayeh A (2023) Leveraging computational psychometrics for language testing. *Language Teaching Research Quarterly* 37: 204–212.

Gil de Zúñiga H and Cheng Z (2021) Origin and evolution of the news finds me perception: Review of theory and effects. *El Profesional de la Información* 30(3): 1–17.

Gil de Zúñiga H, Strauss N and Huber B (2020) The proliferation of the "news finds me" perception across societies. *International Journal of Communication* 14: 1605–1633.

Gil de Zúñiga H, Weeks B and Ardèvol-Abreu A (2017) Effects of the news-finds-me perception in communication: Social media use implications for news seeking and learning about politics. *Journal of Computer-Mediated Communication* 22(3): 105–123.

Goyanes M, Ardèvol-Abreu A and de Zúñiga HG (2021) Antecedents of news avoidance: Competing effects of political interest, news overload, trust in news media, and "news finds me" perception. *Digital Journalism* 11(1): 1–18.

Guzman AL (ed.) (2018) *Human-Machine Communication: Rethinking Communication, Technology, and Ourselves*. New York: Peter Lang.

Imran F and Ali F (2025) Exploring the antecedents of news avoidance behavior among university students in Pakistan. *World of Media* 4: 46–69.

Lewis SC, Guzman AL and Schmidt TR (2019) Automation, journalism, and human–machine communication: Rethinking roles and relationships of humans and machines in news. *Digital Journalism* 7(2): 1–19.

Li Y, Lin X, Sha Z, Jin Z and Li X (2026) AI psychometrics: Evaluating the psychological reasoning of large language models with psychometric validities. In: *Proceedings of the 58th Hawaii International Conference on System Sciences (HICSS)*, pp. 5189–5197.

Maharjan J, Jin R, Zhu J and Kenne D (2025) Psychometric evaluation of large language model embeddings for personality trait prediction. *Journal of Medical Internet Research* 27: e75347.

Markov Č and Min Y (2022) Understanding the public's animosity toward news media: Cynicism and distrust as related but distinct negative media perceptions. *Journalism & Mass Communication Quarterly* 99(4): 1099–1125.

Markov Č and Min Y (2023) Unpacking public animosity toward professional journalism: A qualitative analysis of the differences between media distrust and cynicism. *Journalism* 24(10): 2136–2154.

Nanz A and Matthes J (2022) Democratic consequences of incidental exposure to political information: A meta-analysis. *Journal of Communication* 72(3): 345–373.

Nass C, Steuer J and Tauber ER (1994) Computers are social actors. In: *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '94)*, Boston, MA, 24–28 April 1994, pp. 72–78. New York, NY: Association for Computing Machinery.

Strömbäck J, Tsfati Y, Boomgaarden H, Damstra A, Lindgren E, Vliegenthart R and Lindholm T (2020) News media trust and its impact on media use: Toward a framework for future research. *Annals of the International Communication Association* 44(2): 139–156.

Stroud NJ, Scacco JM and Kim Y (2022) Passive learning and incidental exposure to news. *Journal of Communication* 72(4): 451–460.

Tandoc EC Jr and Kim HK (2023) Avoiding real news, believing in fake news? Investigating pathways from information overload to misbelief. *Journalism* 24(6): 1174–1192.

Tse HTF, Lee FLF and Yang SZQ (2025) Entering journalism in times of democratic backsliding: Hong Kong young journalists' career decision and persistence. *The International Journal of Press/Politics* (Advance online publication).

Tsfati Y, Stromback J, Lindgren E, Boomgaarden HG and Vliegenthart R (2023) What news outlets do people have in mind when they answer survey questions about trust in "media?" *International Journal of Public Opinion Research* 35(2).

Tsfati Y, Vliegenthart R, Stromback J and Lindgren E (2025) An asymmetrical reinforcing spiral? Disentangling the longitudinal dynamics of media use and mainstream media trust. *Journal of Communication* 75(1): 16–26.

Ye H, Huo T, Guo Z, Zhang X, Dong H and Song G (2025) Psychometrics with AI foundation models. *Psychology & Psychiatry Journal* 331.

AI Usage

Anthropic. (2026). Claude Sonnet 4.6 [Large language model]. https://claude.ai
Google. (2026). Gemini 3 Pro [Large language model]. https://gemini.google.com
OpenAI. (2026). ChatGPT 5.3 [Large language model]. https://chatgpt.com

<Used as supplementary tools to verify the subtle nuances and texture of English expressions. The researcher cross-checked the AIs' suggestions and made the final decisions on the text.>